

NAME : PUNDLA ROOPA SUNDAR

REG NO: 21BEC7152

MAIL ID :roopasundar.21bec7152@vitapstudent.ac.in

## IMPORT LIBRARIES

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

## IMPORT DATASET

```
In [3]: df=pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
In [4]: df
```

```
Out[4]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	E
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	
...	...	...	...	...	...	...	...	...	
1465	36	No	Travel_Frequently	884	Research & Development	23	2	Medical	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	Medical	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	Life Sciences	
1468	49	No	Travel_Frequently	1023	Sales	2	3	Medical	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	Medical	

1470 rows × 35 columns

```
In [5]: df.head()
```

Out [5]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	Emp
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	

5 rows × 35 columns

In [6]:

df.tail()

Out [6]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	E
1465	36	No	Travel_Frequently	884	Research & Development	23	2	Medical	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	Medical	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	Life Sciences	
1468	49	No	Travel_Frequently	1023	Sales	2	3	Medical	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	Medical	

5 rows × 35 columns

In [7]:

df.shape

Out [7]:

(1470, 35)

In [8]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object
22	OverTime	1470 non-null	object
23	PercentSalaryHike	1470 non-null	int64
24	PerformanceRating	1470 non-null	int64
25	RelationshipSatisfaction	1470 non-null	int64
26	StandardHours	1470 non-null	int64
27	StockOptionLevel	1470 non-null	int64
28	TotalWorkingYears	1470 non-null	int64
29	TrainingTimesLastYear	1470 non-null	int64
30	WorkLifeBalance	1470 non-null	int64
31	YearsAtCompany	1470 non-null	int64
32	YearsInCurrentRole	1470 non-null	int64
33	YearsSinceLastPromotion	1470 non-null	int64
34	YearsWithCurrManager	1470 non-null	int64

```
dtypes: int64(26), object(9)
```

```
memory usage: 402.1+ KB
```

```
In [9]: df.describe()
```

```
Out[9]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Environm
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	

8 rows × 26 columns

```
In [10]: corr=df.corr()  
corr
```

C:\Users\roopa\AppData\Local\Temp\ipykernel\_22912\3182140910.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.  
corr=df.corr()

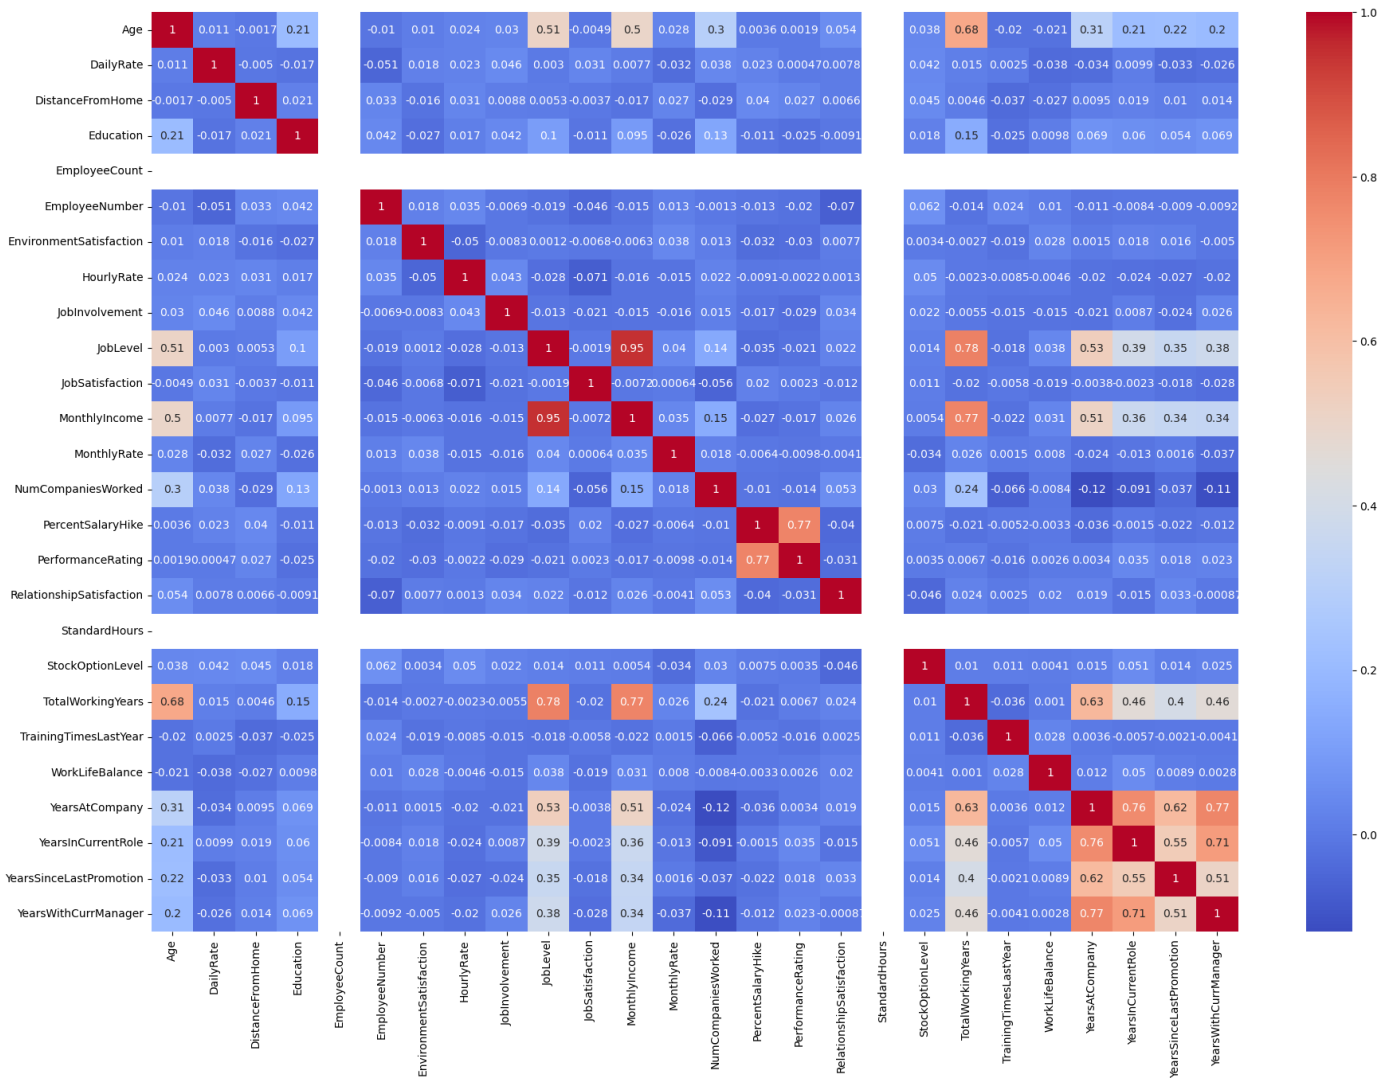
Out[10]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber
Age	1.000000	0.010661	-0.001686	0.208034	NaN	-0.01014
DailyRate	0.010661	1.000000	-0.004985	-0.016806	NaN	-0.05099
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	NaN	0.03291
Education	0.208034	-0.016806	0.021042	1.000000	NaN	0.04207
EmployeeCount	NaN	NaN	NaN	NaN	NaN	NaN
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	NaN	1.00000
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	NaN	0.01762
HourlyRate	0.024287	0.023381	0.031131	0.016775	NaN	0.03517
JobInvolvement	0.029820	0.046135	0.008783	0.042438	NaN	-0.00688
JobLevel	0.509604	0.002966	0.005303	0.101589	NaN	-0.01851
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	NaN	-0.04624
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	NaN	-0.01482
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	NaN	0.01264
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	NaN	-0.00125
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	NaN	-0.01294
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	NaN	-0.02035
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	NaN	-0.06986
StandardHours	NaN	NaN	NaN	NaN	NaN	NaN
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	NaN	0.06222
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	NaN	-0.01436
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	NaN	0.02360
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	NaN	0.01030
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	NaN	-0.01124
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	NaN	-0.00841
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	NaN	-0.00901
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	NaN	-0.00919

26 rows × 6 columns

```
In [11]: plt.subplots(figsize=(22,15))  
sns.heatmap(corr,annot=True,cmap="coolwarm")
```

Out[11]: <Axes: >



```
In [12]: df.Attrition.value_counts()
```

```
Out[12]: No      1233
         Yes      237
         Name: Attrition, dtype: int64
```

Checking for NULL Values

```
In [13]: df.isnull().any()
```


```
Out[13]: Age False
Attrition False
BusinessTravel False
DailyRate False
Department False
DistanceFromHome False
Education False
EducationField False
EmployeeCount False
EmployeeNumber False
EnvironmentSatisfaction False
Gender False
HourlyRate False
JobInvolvement False
JobLevel False
JobRole False
JobSatisfaction False
MaritalStatus False
MonthlyIncome False
MonthlyRate False
NumCompaniesWorked False
Over18 False
OverTime False
PercentSalaryHike False
PerformanceRating False
RelationshipSatisfaction False
StandardHours False
StockOptionLevel False
TotalWorkingYears False
TrainingTimesLastYear False
WorkLifeBalance False
YearsAtCompany False
YearsInCurrentRole False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool
```

## Data Visualization

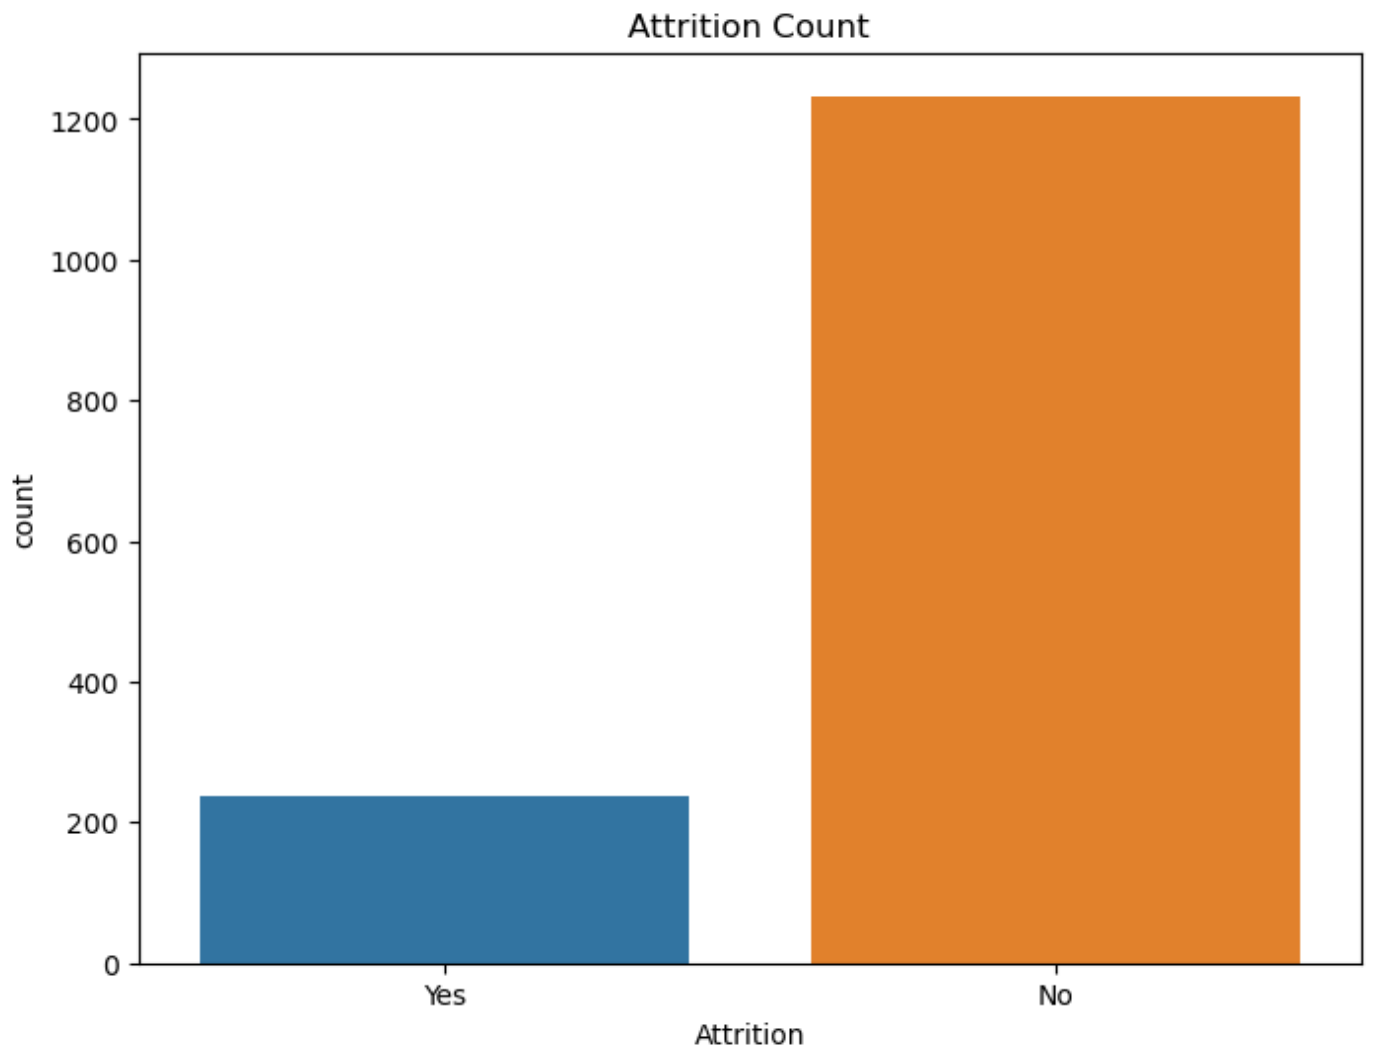
```
In [14]: attrition_counts = df['Attrition'].value_counts()
plt.figure(figsize=(6, 6))
plt.pie(attrition_counts, labels=attrition_counts.index, autopct='%1.1f%%', startangle=90)
plt.title('Attrition Distribution')
plt.axis('equal')

plt.show()
```

Attrition Distribution




```
In [15]: plt.figure(figsize=(8, 6))  
sns.countplot(x="Attrition", data=df)  
plt.title("Attrition Count")  
plt.show()
```




```
In [16]: plt.figure(figsize=(8, 6))
sns.histplot(data=df, x="Age", kde=True)
plt.title("Distribution of Age")
plt.show()
```






### Outlier Detection

```
In [17]: plt.figure(figsize=(35, 8))
sns.boxplot(data=df)
plt.title('Box Plots for all the attributes')
plt.show()
```




```
In [18]: sns.boxplot(data=df, x='YearsInCurrentRole')
plt.title('Years In Current Role')
plt.xlabel('Fare')
plt.show()
```

Years In Current Role



```
In [19]: sns.boxplot(data=df, x='MonthlyIncome')
plt.title('Monthly Income')
plt.xlabel('Fare')
plt.show()
```


Monthly Income



```
In [20]: from scipy import stats
```

```
z_scores = stats.zscore(df['MonthlyIncome'])
z_score_threshold = 3
df_cleaned = df[(np.abs(z_scores) <= z_score_threshold)]
```

```
In [21]: sns.boxplot(data=df_cleaned, x='MonthlyIncome')
plt.title('Monthly Income')
plt.xlabel('Fare')
plt.show()
```



So the outliers are in large quantity, and they are inside the threshold, so let us not remove the outliers

## SPLITTING INDEPENDENT AND DEPENDENT VARIABLES

```
In [22]: x = df.drop(columns=["Attrition"])
y = df["Attrition"]
```

```
In [23]: x.head()
```

```
Out[23]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCour
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	

5 rows × 34 columns

```
In [24]: y.head()
```

```
Out[24]: 0    Yes
1    No
2    Yes
3    No
4    No
Name: Attrition, dtype: object
```

## ENCODING

```
In [25]: categorical_features = x.select_dtypes(include=['object']).columns.tolist()
x_encoded = pd.get_dummies(x, columns=categorical_features, drop_first=True)
```

```
In [26]: x_encoded.head()
```

```
Out[26]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	...
0	41	1102	1	2	1	1	2	...
1	49	279	8	1	1	2	3	...
2	37	1373	2	2	1	4	4	...
3	33	1392	3	4	1	5	4	...
4	27	591	2	1	1	7	1	...

5 rows × 47 columns

## FEATURE SCALING

```
In [27]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
x_scaled = pd.DataFrame(scaler.fit_transform(x_encoded), columns=x_encoded.columns)
```

```
In [28]: x_scaled.head()
```

```
Out[28]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	...
0	0.446350	0.742527	-1.010909	-0.891688	0.0	-1.701283	-0.6609	...
1	1.322365	-1.297775	-0.147150	-1.868426	0.0	-1.699621	0.2546	...
2	0.008343	1.414363	-0.887515	-0.891688	0.0	-1.696298	1.1697	...
3	-0.429664	1.461466	-0.764121	1.061787	0.0	-1.694636	1.1697	...
4	-1.086676	-0.524295	-0.887515	-1.868426	0.0	-1.691313	-1.5756	...

5 rows × 47 columns

```
In [29]: x=x_scaled
```

Train and test split

```
In [30]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [31]: # Import the necessary libraries
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from joblib import dump
```

```
In [32]: logreg_model = LogisticRegression(random_state=42)
dt_model = DecisionTreeClassifier(random_state=42)
```

```
In [33]: logreg_model.fit(x_train, y_train)
dt_model.fit(x_train, y_train)
```

```
Out[33]: ▼      DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

```
In [34]: logreg_predictions = logreg_model.predict(x_test)

dt_predictions = dt_model.predict(x_test)

logreg_accuracy = accuracy_score(y_test, logreg_predictions)
print("Logistic Regression Accuracy:", logreg_accuracy)

dt_accuracy = accuracy_score(y_test, dt_predictions)
print("Decision Tree Accuracy:", dt_accuracy)

logreg_report = classification_report(y_test, logreg_predictions)
print("Classification Report for Logistic Regression:\n", logreg_report)

dt_report = classification_report(y_test, dt_predictions)
print("Classification Report for Decision Tree Classifier:\n", dt_report)

logreg_conf_matrix = confusion_matrix(y_test, logreg_predictions)
print("Confusion Matrix for Logistic Regression:\n", logreg_conf_matrix)

dt_conf_matrix = confusion_matrix(y_test, dt_predictions)
print("Confusion Matrix for Decision Tree Classifier:\n", dt_conf_matrix)
```

Logistic Regression Accuracy: 0.8809523809523809  
Decision Tree Accuracy: 0.7721088435374149  
Classification Report for Logistic Regression:

	precision	recall	f1-score	support
No	0.92	0.95	0.93	255
Yes	0.56	0.46	0.51	39
accuracy			0.88	294
macro avg	0.74	0.70	0.72	294
weighted avg	0.87	0.88	0.88	294

Classification Report for Decision Tree Classifier:

	precision	recall	f1-score	support
No	0.87	0.86	0.87	255
Yes	0.17	0.18	0.17	39
accuracy			0.77	294
macro avg	0.52	0.52	0.52	294
weighted avg	0.78	0.77	0.78	294

Confusion Matrix for Logistic Regression:  
[[241 14]  
[ 21 18]]  
Confusion Matrix for Decision Tree Classifier:  
[[220 35]  
[ 32 7]]

In [ ]:

In [ ]:

In [ ]: